

# Supervised Learning with scikit-learn

## 4). Preprocessing & Pipelining:

### a). Exploring Categorical Features

```
# Import pandas
```

```
import pandas as pd
```

```
# Read 'gapminder.csv' into a DataFrame: df
```


```
df = pd.read_csv('gapminder.csv')
```

```
# Create a boxplot of life expectancy per region
```

```
df.boxplot('life', 'Region', rot=60)
```

```
# Show the plot
```

```
plt.show()
```



**b). Creating Dummy Variables**

```
# Create dummy variables: df_region
```

```
df_region = pd.get_dummies(df)
```

```
# Print the columns of df_region
```

```
print(df_region.columns)
```

```
# Drop 'Region_America' from df_region
```

```
df_region = pd.get_dummies(df, drop_first=True)
```

```
# Print the new columns of df_region
```

```
print(df_region.columns)
```

<script.py> output:

```
Index(['population', 'fertility', 'HIV', 'CO2', 'BMI_male', 'GDP',  
      'BMI_female', 'life', 'child_mortality', 'Region_America',  
      'Region_East Asia & Pacific', 'Region_Europe & Central Asia',  
      'Region_Middle East & North Africa', 'Region_South Asia',  
      'Region_Sub-Saharan Africa'],  
      dtype='object')
```

```
Index(['population', 'fertility', 'HIV', 'CO2', 'BMI_male', 'GDP',  
      'BMI_female', 'life', 'child_mortality', 'Region_East Asia & Pacific',  
      'Region_Europe & Central Asia', 'Region_Middle East & North Africa',  
      'Region_South Asia', 'Region_Sub-Saharan Africa'],  
      dtype='object')
```

**c). Regression with categorical features****# Import necessary modules****from sklearn.linear\_model import Ridge****from sklearn.model\_selection import cross\_val\_score****# Instantiate a ridge regressor: ridge****ridge = Ridge(alpha=0.5, normalize=True)****# Perform 5-fold cross-validation: ridge\_cv****ridge\_cv = cross\_val\_score(ridge,X,y,cv=5)****# Print the cross-validated scores****print(ridge\_cv)**

&lt;script.py&gt; output:

**[ 0.86808336 0.80623545 0.84004203 0.7754344 0.87503712]**

**d). Dropping Missing Data:**

```
# Convert '?' to NaN
```

```
df[df == '?'] = np.nan
```

```
# Print the number of NaNs
```

```
print(df.isnull().sum())
```

```
# Print shape of original DataFrame
```

```
print("Shape of Original DataFrame: {}".format(df.shape))
```

```
# Drop missing values and print shape of new DataFrame
```

```
df = df.dropna()
```

```
# Print shape of new DataFrame
```

```
print("Shape of DataFrame After Dropping All Rows with Missing Values: {}".format(df.shape))
```

<script.py> output:

|                   |     |
|-------------------|-----|
| party             | 0   |
| infants           | 12  |
| water             | 48  |
| budget            | 11  |
| physician         | 11  |
| salvador          | 15  |
| religious         | 11  |
| satellite         | 14  |
| aid               | 15  |
| missile           | 22  |
| immigration       | 7   |
| synfuels          | 21  |
| education         | 31  |
| superfund         | 25  |
| crime             | 17  |
| duty_free_exports | 28  |
| eea_rsa           | 104 |

dtype: int64

Shape of Original DataFrame: (435, 17)

Shape of DataFrame After Dropping All Rows with Missing Values: (232, 17)

**e). Imputing missing data in a ML Pipeline****# Import necessary modules****from sklearn.preprocessing import Imputer****from sklearn.pipeline import Pipeline****from sklearn.svm import SVC****# Setup the pipeline steps: steps****steps = [('imputation', Imputer(missing\_values='NaN', strategy='most\_frequent', axis=0)),  
('SVM', SVC())]****# Create the pipeline: pipeline****pipeline = Pipeline(steps)****# Create training and test sets****X\_train, X\_test, y\_train, y\_test = train\_test\_split(X,y,test\_size=0.3,random\_state=42)****# Fit the pipeline to the train set****pipeline.fit(X\_train,y\_train)****# Predict the labels of the test set****y\_pred = pipeline.predict(X\_test)****# Compute metrics****print(classification\_report(y\_test, y\_pred))**

&lt;script.py&gt; output:

precision recall f1-score support

democrat 0.99 0.96 0.98 85

republican 0.94 0.98 0.96 46

avg / total 0.97 0.97 0.97 131

**f). Centring and Scaling your Data****# Import scale****from sklearn.preprocessing import scale****# Scale the features: X\_scaled****X\_scaled = scale(X)****# Print the mean and standard deviation of the unscaled features****print("Mean of Unscaled Features: {}".format(np.mean(X)))****print("Standard Deviation of Unscaled Features: {}".format(np.std(X)))****# Print the mean and standard deviation of the scaled features****print("Mean of Scaled Features: {}".format(np.mean(X\_scaled)))****print("Standard Deviation of Scaled Features: {}".format(np.std(X\_scaled)))**

&lt;script.py&gt; output:

Mean of Unscaled Features: 18.432687072460002

Standard Deviation of Unscaled Features: 41.54494764094571

Mean of Scaled Features: 2.7314972981668206e-15

Standard Deviation of Scaled Features: 0.9999999999999999

**g). Centering and Scaling in a Pipeline**

**# Import the necessary modules**

**from sklearn.preprocessing import StandardScaler**

**from sklearn.pipeline import Pipeline**

**# Setup the pipeline steps: steps**

```
steps = [('scaler', StandardScaler()),  
        ('knn', KNeighborsClassifier())]
```

**# Create the pipeline: pipeline**

```
pipeline = Pipeline(steps)
```

**# Create train and test sets**

```
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3, random_state=42)
```

**# Fit the pipeline to the training set: knn\_scaled**

```
knn_scaled = pipeline.fit(X_train,y_train)
```

**# Instantiate and fit a k-NN classifier to the unscaled data**

```
knn_unscaled = KNeighborsClassifier().fit(X_train, y_train)
```

**# Compute and print metrics**

```
print('Accuracy with Scaling: {}'.format(knn_scaled.score(X_test, y_test)))
```

```
print('Accuracy without Scaling: {}'.format(knn_unscaled.score(X_test, y_test)))
```

**<script.py> output:**

Accuracy with Scaling: 0.7700680272108843

Accuracy without Scaling: 0.6979591836734694



**h). Pipeline for Classification****# Setup the pipeline**

```
steps = [('scaler', StandardScaler()),
        ('SVM', SVC())]
```

**pipeline = Pipeline(steps)****# Specify the hyperparameter space**

```
parameters = {'SVM__C':[1, 10, 100],
              'SVM__gamma':[0.1, 0.01]}
```

**# Create train and test sets**

```
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2,random_state=21)
```

**# Instantiate the GridSearchCV object: cv**

```
cv = GridSearchCV(pipeline,param_grid=parameters)
```

**# Fit to the training set**

```
cv.fit(X_train,y_train)
```

**# Predict the labels of the test set: y\_pred**

```
y_pred = cv.predict(X_test)
```

**# Compute and print metrics**

```
print("Accuracy: {}".format(cv.score(X_test, y_test)))
```

```
print(classification_report(y_test, y_pred))
```

```
print("Tuned Model Parameters: {}".format(cv.best_params_))
```

&lt;script.py&gt; output:

```
Accuracy: 0.7795918367346939
```

```

precision  recall  f1-score  support
False      0.83    0.85    0.84    662
True       0.67    0.63    0.65    318
avg / total  0.78    0.78    0.78    980
```

```
Tuned Model Parameters: {'SVM__gamma': 0.1, 'SVM__C': 10}
```

**i). Pipelining for Regression**

**# Setup the pipeline steps: steps**

```
steps = [('imputation', Imputer(missing_values='NaN', strategy='mean', axis=0)),  
        ('scaler', StandardScaler()),  
        ('elasticnet', ElasticNet())]
```

**# Create the pipeline: pipeline**

```
pipeline = Pipeline(steps)
```

**# Specify the hyperparameter space**

```
parameters = {'elasticnet__l1_ratio': np.linspace(0,1,30)}
```

**# Create train and test sets**

```
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.4, random_state=42)
```

**# Create the GridSearchCV object: gm\_cv**

```
gm_cv = GridSearchCV(pipeline, param_grid=parameters)
```

**# Fit to the training set**

```
gm_cv.fit(X_train,y_train)
```

**#Compute and print the metrics**

```
r2 = gm_cv.score(X_test, y_test)
```

```
print("Tuned ElasticNet Alpha: {}".format(gm_cv.best_params_))
```

```
print("Tuned ElasticNet R squared: {}".format(r2))
```

<script.py> output:

```
Tuned ElasticNet Alpha: {'elasticnet__l1_ratio': 1.0}
```

```
Tuned ElasticNet R squared: 0.8862016570888217
```